

Please use Piazza:

<https://piazza.com/class/mk7c83n7ab92h3>

<http://dsg.csail.mit.edu/6.5830/>

6.5830/6.5831

Database Systems

6.5830 Lecture 1- 2/3/2026


Mike Cafarella, Tianyu Li

michjc@csail.mit.edu, litianyu@mit.edu

Michael Cafarella




Tianyu Li



Eugenie Lai



Tamar Korkotashvili



Ivy Zheng

Administrivia

<http://dsg.csail.mit.edu/6.5830>, Email: 6.5830-staff@mit.edu

Ask questions on Piazza! You should have already been added.

Lecturers:

Mike Cafarella (michjc@mit.edu)

Tianyu Li (litianyu@mit.edu)

TAs:

Tamar Korkotashvili (tkorkot@mit.edu)

Ivy Zheng (ivyzheng@mit.edu)

Eugenie Lai (eylai@mit.edu)

Office hours:

see website + piazza announcement for changes


What is a Database?

- Structured Data Collection
 - Records
 - Relationships
- This class: Database Management Systems (DBMSs)

Software systems for storing and querying databases

Why should I care?

1960s



IDS



Mainframes



Enterprise
Record
Processing

1970s-80s



System R,
Ingres, etc.



Early
servers



SQL,
Transactions

2000s



NoSQL
Systems



Commodity
Clusters



"Web-scale"

2015-2025



Cloud DBs




Public
Cloud



Elastic,
Real-time
Processing
& Analytics

2025- ?




DB for AI



GPU/TPU
Clusters



Agentic
Workloads



Course Overview

- This class focuses on the themes in data management systems that will be relevant long after I am dead (probably)
- Focus on classical relational system, but will branch out later in the semester
- Readings are drawn from literature and posted on course website

6.5830/1 Concepts

- Data modeling / layout
 - Declarative querying
 - Query processing
 - Algorithms for accessing and manipulating data
-
- Consistency / Transactions (“ACID”)
 - “Big Data” – scaling to massive volumes, many machines
 - Semantic Data Management – managing and querying semantic information in unstructured data

Quiz 1

Quiz 2

Two class flavors

6.5831 is an undergraduate class designed to satisfy the **AUS requirement** in the EECS curriculum (instead of 6.1800). The class does **not** fulfill the CI-M requirement.

6.5830 is a Grad-H class. It counts as an engineering concentration (EC) subject in Systems. **For Area II Ph.D. students in EECS, it satisfies the Systems TQE requirement and the AUS requirement.** 6.5830 requires the completion of a final open-ended research project

Pre-requisites

- 6.1210 or 6.1220 (or equivalent)
 - i.e., introduction to algorithms
- 6.1800 or 6.1810 (or equivalent)
 - i.e., introduction to computer systems
- Courage and grit to do low-level systems programming
 - You should want to do bit-twiddling, memory management, concurrency

6.5830/1 Assignments

- 3 Labs: GoDB
- 3 Problem Sets
- 2 Quizzes
 - All on gradescope. You should have been added.
- 6.5830: **Final Project** – open ended project of your choice
- 6.5831: Lab 4 **OR** open-ended project

GoDB (v2)

- New version of educational system this semester
 - Designed to teach more real-world system programming practices
 - More emphasis on performance, concurrency, algorithm design
- ALWAYS start early on assignments
 - Brush up on system programming basics and come to OH if you need help

Open-ended Final Project

- You can propose your own data systems research
 - Subject to instructor's discretion
- New this semester: OSS Partners
 - Curated projects from Apache Data Fusion, ChromaDB, Neon, and Palimpzest/Carnot
- New this semester: Practicum
 - Extend GoDB with a non-trivial feature
- Final deliverable: 5-min recorded presentation, in-class Q&A, written report

AI Policy

- You are allowed to use generative AI tools
 - For labs: however you want
 - For psets: for background help; no pasting of questions into context
- You must use tools that are equally accessible to everyone in the class
 - E.g., free-tier ChatGPT, free Gemini Pro for students

AI Policy (cont.)

- You are responsible for your AI usage
 - Course staff will not help you debug vibe code, nor will you get points back if the AI misled you
- **New:** Lab Write-up will happen in class
 - Usually a week after labs are due
 - Think of them as mini quizzes to test your understanding of your code and the lab content

Collaboration Policy


- Psets are individual assignments
- You are allowed to complete labs in pairs
 - Write-ups are individually completed and graded
- Sharing answers/code publicly or between teams is not allowed
 - Discussion of general ideas is allowed
 - Collaboration violation yields a zero for all parties, and possibly referral to committee on discipline, etc.
 - Please don't do it!

Collaboration Policy (cont.)

- You are allowed to complete final projects in teams of 3
 - Team matching is enabled on piazza
 - Submit project teams by 2/26
- You are expected to have the same lab & project teams throughout the semester
 - Consult course staff if you need to change team

Resources

- In-person office hours posted on the website
- Use Piazza for conceptual and practical questions
 - Please don't post code
 - We try to respond in a timely manner, but no guarantees
- Email or private Piazza messages are fine for personal matters



Lab 1 Out Today: Storage & Buffer Management
PS1 Out Today: Designed to teach you SQL

Today

- Why database systems?
- User's perspective:
 - Modeling data
 - Querying data
- Data Models

Zoo Website Features


- Admin interface
 - Edit
 - Add an animal
- Public
 - Pictures & Maps
- Zookeeper
 - Feed times



- 1K animals, 5K URLs, 10 admins, 200 keepers

Zoo Data Model

Entity Relationship Diagram



Animals have names, ages, species

Keepers have names

Cages have cleaning times, buildings

Animals are in 1 cage; cages have multiple animals

Keepers keep multiple cages, cages kept by multiple keepers

Our Zoo



Mike the Moose




Tianyu the Giraffe



Sally the Student


Zoo Data Model

Entity Relationship Diagram



Zoo Data Model

Entity Relationship Diagram



Cageid	Feedtime	Bldg
1	1:30	1
2	2:30	2

keeperid	name
1	jenny
2	joe

keeperid	cageid
1	1
1	2
2	1

Study Break #1

- Questions
 - Are there other ways to represent this zoo data than a collection of tables?
 - What are tradeoffs in different representations?


Alternatives to Relations

Hierarchy

cage 1
 tianyuan
 giraffe
 13 yrs
 sally
 student
 1 yr

cage 2
 mike
 moose
 3 yrs

Graph



Multiple Tabular Representations Are Possible

name	age	species	cageno	feedtime	bldg
tianyu	13	giraffe	1	1:30	1
mike	3	moose	2	2:30	2
sally	1	student	1	1:30	1

Is this a good representation? Why or why not?

Not “Normalized” – repeats data. More in later lectures!

SQL – Structured Query Language

```
SELECT  field1, ..., fieldM  
FROM    table1, ...  
WHERE   condition1, ...
```

```
INSERT INTO table VALUES (field1, ...)
```

```
UPDATE table SET field1 = X, ...  
WHERE condition1,...
```

Names of Giraffes

- Imperative

```
for each row r in animals
    if r.species = 'giraffe'
        output r.name
```

- Declarative

```
SELECT r.name FROM animals
WHERE r.species = 'giraffe'
```

Cages in Building 32

**NESTED
LOOPS**

- Imperative

```
for each row a in animals
  for each row c in cages
    if a.cageno = c.no and c.bldg = 32
      output a
```

- Declarative

JOIN

```
SELECT a.name FROM animals AS a, cages AS c
WHERE a.cageno = c.no AND c.bldg = 32
```


Average Age of Bears

- Declarative

```
SELECT AVG(age) FROM animals  
WHERE species = 'bear'
```

Complex Queries

Find pairs of animals of the same species and different genders older than 1 year:

```
SELECT a1.name,a2.name  
FROM animals as a1, animals as a2  
WHERE a1.gender = M and a2.gender = F  
AND a1.species = a2.species  
AND a1.age > 1 and a2.age > 1
```

“self join”

Find cages with salamanders fed later than the average feedtime of any cage:

```
SELECT cages.cageid FROM cages, animals  
WHERE animals.species = 'salamander'  
AND animals.cageid = cages.cageid  
AND cages.feedtime >  
    (SELECT AVG(feedtime) FROM cages )
```

“nested queries”

Complex Queries 2

Find keepers who keep both students and salamanders:

```
SELECT keeper.name
```

Complex Queries 2

Find keepers who keep both students and salamanders:

```
SELECT keeper.name
```

```
FROM keeper, cages as c1, cages as c2,
```

```
     keeps as k1, keeps as k2, animals as a1, animals as a2
```

Complex Queries 2

Find keepers who keep both students and salamanders:

```
SELECT keeper.name  
FROM keeper, cages as c1, cages as c2,  
      keeps as k1, keeps as k2, animals as a1, animals as a2  
WHERE c1.cageid = k1.cageid AND keeper.keeperid = k1.keeperid  
AND c2.cageid = k2.cageid AND keeper.keeperid = k2.keeperid
```

Complex Queries 2

Find keepers who keep both students and salamanders:

```
SELECT keeper.name
```

```
FROM keeper, cages as c1, cages as c2,
```


```
    keeps as k1, keeps as k2, animals as a1, animals as a2
```

```
WHERE c1.cageid = k1.cageid AND keeper.keeperid = k1.keeperid
```

```
AND c2.cageid = k2.cageid AND keeper.keeperid = k2.keeperid
```

```
AND a1.species = 'student' AND a2.species = 'salamander'
```

```
AND c1.cageid = a1.cageid AND c2.cageid = a2.cageid
```




Declarative Queries: **What**, not **How**


- Many possible procedures to answer a query
- Besides looping through the table, what else could we do?
 - Sort animals on type
 - + good for “bears” query
 - Inserts are slower
 - Store animals table in a hash table or tree (“index”)



SQL → Procedural Plan → Optimized
Plan → Compiled Program



SQL → Procedural Plan → Optimized Plan → Compiled Program



SQL programmer just thinks in terms of table operations, not the order or implementation!

Summary: Database Systems

- Relational Model + Schema Design
- Declarative Queries
- Query Optimization
- Efficient access and updates to data
 - Recoverability
 - Consistency

Relational Model


“Those who cannot remember the past are doomed to repeat it”

A Short History Lesson

- **Different Data Models**
 - Hierarchical (IMS/DL1) – 1960's
 - Network (CODASYL) – 1970's
 - Relational – 1970's and beyond
- **Key ideas**
 - Data redundancy (and how to avoid it)
 - Physical and logical data independence
 - Relational algebra and axioms

Recap: Zoo Data Model

Entity Relationship Diagram



Animals have names, ages, species

Keepers have names

Cages have cleaning times, buildings

Animals are in 1 cage; cages have multiple animals

Keepers keep multiple cages, cages kept by multiple keepers

Zoo Tables (aka Relations)

Animals

id	name	age	species	cageno
1	Mike	3	Moose	1
2	Tianyu	12	Giraffe	1
3	Sally	1	Student	2

“Schema”: Field names
& types

Rows, records, or tuples

Cages

no	feedtime	building
1	12:30	1
2	1:30	2


Keepers

id	name
1	Jane
2	Joe

Keeps

kid	cageno
1	1
1	2
2	1

Modified Zoo Data Model



Slightly different compared to last time:

- Each animal in 1 cage, multiple animals share a cage
- Each animal cared for by 1 keeper, keepers care for multiple animals


IMS (Hierarchical Model)

- Data organized as *segments*
 - Collection of records, each with same *segment type*
 - Arranged in a tree of segment types, e.g.:



- Segments have different physical representations
 - Unordered
 - Indexed
 - Sorted
 - Hashed

Example Hierarchy



IMS Physical Representation

Keepers segment

A1 Segment

C1 Segment

A2 Segment

C2 Segment

A3 Segment

C3 Segment

Segment Structure

- Each segment has a particular physical representation
 - Chosen by database administrator
 - E.g., ordered, hashed, unordered...
- Choice of segment structure affects which operations can be applied on it

IMS / DL/1 Operations

- **GetUnique** (seg type, pred)
 - Get first record satisfying pred
 - Only supported by hash / sorted segments
- **GetNext** (seg type, pred)
 - Get first or next key in hierarchical order
 - Starts from last GetNext/GetUnique call
- **GetNextParent** (seg type, pred)
 - Same as GetNext, but will not move up hierarchy to next parent
- **Delete, Insert**

Example PL/1 Program #1

Find the cages that Jane keeps

```
GetUnique(Keepers, name = "Jane")
```

Until done:

```
    cageid = GetNextParent (cages).no  
    print cageid
```

Jane (keeper) (HSK 1)

Mike, moose, ...	(2)
1, 100sq ft, ...	(3)
Tianyu, giraffe, ...	(4)
2, 1000sq ft, ...	(5)
Sally, student, ...	(6)
1, 100sq ft, ...	(7)

Joe (keeper) (8)

This iterates through
data underneath
Jane

Implicitly, now
navigating from the
Jane record in
keepers

Example PL/1 Program #2

Find the keepers that keep cage 6

```
keep = GetUnique(keepers)
```

Until done:

```
    cage = GetNextParent(cages, id = 6)
```

```
    if (cage is not null):
```

```
        print keep
```

```
    keep = GetNext(keepers)
```

What's Bad About IMS/PL1?

- Duplication of data w/ non-hierarchical data
- Painful low level programming interface – have to program the search algorithm
- Limited **physical data independence**
 - Change root from indexed to hash --- programs that do GetNext on the root segment will fail
 - Change root from keepers to animals? Also fails.
 - Cannot do inserts into sequential root structure
- Limited **logical data independence**
 - Schemas change, do programs have to?

Logical Data Independence

- Suppose as a cost cutting measure, Zoo management decides a keeper will be responsible for a cage – and all the animals in that cage.




Programs have to change, because the position in the database after a GN/GNP call may not be the same anymore!

Will see how SQL addresses this


Schemas Change for Many Reasons

- Management decides to have “patrons” who buy cages
 - Need to add a patronid column
- Feds change the rules (OSHA)
 - Keepers can keep at most 2 cages
- Tax rules change (IRS)
- Merge with another zoo



Study break #2

- Consider a course schema with students, classes, rooms (each has a number of attributes)



Classes in exactly one room


Students in zero or more classes

Classes taken by zero or more students

Rooms host zero or more classes

Questions

1. Describe one possible hierarchical schema for this data
2. Is there a hierarchical representation that is free of redundancy?



Solution

- Many are possible; one example:
 - Classes
 - Students
 - Rooms
- Duplicates data about students,
 - Students take multiple classes, rooms host multiple classes
- Any other arrangement also duplicates data

CODASYL



- **Conference/Committee on Data Systems Languages**
 - Responsible for COBOL
- CODASYL data model developed by consortium of large companies in the 70's
- Designed to address limitations of IMS/PL1
- Graph or network-based data model



The Computer Museum in Boston celebrated cobol's 25th anniversary on May 16, 1985. The cobol tombstone sent to Charles Phillips (see his adjoining article) was presented to the museum. Here surrounding the tombstone (left to right): Ron Hamm, Jack Jones, Jan Prokop, Oliver Smoot, Tom

Rice, Donald Nelson, Grace Hopper, Michael O'Connell, and Howard Bromberg (photo by Lilian Kemp). At the museum's celebration, Bromberg told the following tale of the tombstone.

Example CODASYL Network



Records can either be hashes (allowing equality lookup) or sorted ("clustered") according to some key (allowing a range lookup).


Example: Find Cages Joe Keeps

Find keepers (name = 'Joe')

Until done:

Find next animal in caredforby

Find cage in livesin



- Programming is finding an entry point and navigating around in multidimensional space
 - Each line of code is implicitly at some location in this structure
 - Have to remember where you are

CodasyI Problems

- Incredibly complex —
“Navigational Programming”
- Programs lack physical or logical data independence
 - Can't change schema w/out changing programs;
 - Can't change physical representation either b/c different index types might or might not support different operations
- Some of this could have been fixed by adding a high-level language to CODASYL
- Relational model was a clean-slate approach designed to fix this



Relational Principles

- Simple representation
- Set-oriented programming model that doesn't require "navigation"
- No physical data model description required(!)
 - E.g., no specification of sort orders, hashes, etc

Relational Data Model

- All data is represented as tables of records (*tuples*)
- Tables are unordered sets (no duplicates)
- Database is one or more tables
- Each relation has a *schema* that describes the types of the columns/fields
- Each field is a primitive type -- not a set or relation
- Physical representation/layout of data is not specified (no index types, nestings, etc)

Zoo Tables

Foreign
Keys

Animals

id	name	age	species	cageno	keptby	feedtime
1	Mike	3	Moose	1	1	10:00 am
2	Tianyu	12	Giraffe	1	2	11:00 am
3	Sally	1	Student	2	1	1:00 pm

Primary
Key

Cages

no	building
1	1
2	2

Primary
Key

Keepers

id	name
1	Jane
2	Joe

Primary
Key

Schema: Animals

(id: int,
name: string,
age: int,
species: string,
cageno: int **references** cages.no,
keptby: int **references** keepers.id.
feedtime: time)

Zoo Tables (original schema)

Animals

id	name	age	species	cageno
1	Mike	3	Moose	1
2	Tianyu	12	Giraffe	1
3	Sally	1	Student	2

Primary
Key

Foreign
Key

Cages

no	feedtime	building
1	12:30	1
2	1:30	2

Primary
Key

Keepers

id	name
1	Jane
2	Joe

Primary
Key

Keeps

kid	cageno
1	1
1	2
2	1

Foreign
Key

Foreign
Key

Relational Algebra

- **Projection** ($\pi(T, c_1, \dots, c_n)$)
 - select a subset of columns $c_1 \dots c_n$
- **Selection** ($\sigma(T, \text{pred})$)
 - select a subset of rows that satisfy pred
- **Cross Product** ($T_1 \times T_2$)
 - combine two tables
- **Join** ($\bowtie(T_1, T_2, \text{pred})$) = $\sigma(T_1 \times T_2, \text{pred})$
 - combine two tables with a predicate
- Plus set operations (UNION, DIFFERENCE, etc)
- “Algebra” – Closed under its own operations
 - Every expression over relations produces a relation

Join as Cross Product

Animals

name	cageno
Mike	1
Tianyu	1
Sally	2

Cages

no	bldg
1	32
2	36

cageno	no	name	bldg
1	1	Mike	32
1	2	Mike	36
1	1	Tianyu	32
1	2	Tianyu	36
2	1	Sally	32
2	2	Sally	36

Find animals in bldg. 32

```
σ (  
  ⋈(  
    animals,  
    cages,  
    animals.cageno = cages.no  
  ),  
  bldg = 32  
)
```

**Real implementations do not ever
materialize the cross product**

Join as Cross Product

Animals

name	cageno
Sam	1
Tianyu	1
Sally	2

Cages

no	bldg
1	32
2	36

cageno	no	name	bldg
1	1	Mike	32
1	2	Mike	36
1	1	Tianyu	32
1	2	Tianyu	36
2	1	Sally	32
2	2	Sally	36

Find animals in bldg. 32

1. animals.cageno = cages.no

```
σ (  
  ⋈ (  
    animals,  
    cages,  
    animals.cageno = cages.no  
  ),  
  bldg = 32  
)
```

Join as Cross Product

Animals

name	cageno
Sam	1
Tianyu	1
Sally	2

Cages

no	bldg
1	32
2	36

cageno	no	name	bldg
1	1	Mike	32
1	2	Mike	36
1	1	Tianyu	32
1	2	Tianyu	36
2	1	Sally	32
2	2	Sally	36

Find animals in bldg. 32

```
σ (  
  ⋈ (  
    animals,  
    cages,  
    animals.cageno = cages.no  
  ),  
  bldg = 32  
)
```

1. animals.cageno = cages.no
2. **bldg = 32**

*Do you think this is how
databases actually
execute joins?*

Relational Identities

- **Join reordering**

- $A \bowtie B = B \bowtie A$
- $(A \bowtie B) \text{ join } C = A \bowtie (B \bowtie C)$

- **Selection reordering**

- $\sigma_1(\sigma_2(A)) = \sigma_2(\sigma_1(A))$

- **Selection push down**


- $\sigma(A \bowtie_{\text{pred}} B) = \sigma(A) \bowtie_{\text{pred}} \sigma(b)$
- σ may only apply to one table

- **Projection push down**

- $\pi(\sigma(A)) = \sigma(\pi(A))$
- As long as π doesn't remove fields used in σ
- Also applies to joins

Push Down Example

```
σ (  
  ⋈ (  
    animals,  
    cages,  
    animals.cageno = cages.no  
  ),  
  bldg = 32  
)
```



```
⋈ (  
  animals,  
  σ (  
    cages,  
    bldg = 32  
  )  
  animals.cageno = cages.no  
)
```

Join Ordering Example

- Find buildings Joe keeps
- SQL

```
SELECT building
FROM cages JOIN keeps ON no = cageno
JOIN keepers on kid = id
WHERE name = 'Joe'
```

*SQL query executor
free to choose either
ordering!*

*Text of SQL query is
not an ordering*

```
⋈ (
  ⋈ (
    cages,
    keeps,
    no = cageno
  ),
  σ (
    keepers,
    name = 'Joe'
  ),
  kid = id
)
```



```
⋈ (
  cages,
  ⋈ (
    σ (
      keepers,
      name = 'Joe'
    ),
    keeps,
    kid=id
  ),
  no = cageno
)
```

*Best
ordering
depends on
sizes of
tables*

*Filtered
keepers may
be **much**
smaller*

Study Break # 2

Schema:

classes: (cid, c_name, c_rid, ...)

rooms: (rid, bldg, ...)

students: (sid, s_name, ...)

takes: (t_sid, t_cid)

```
SELECT s_name FROM student,takes,classes
WHERE t_sid=sid AND t_cid=cid
AND c_name='6.830'
```

- Write an equivalent relational algebra expression for this query
- Are there other possible expressions?
- Do you think one would be more “efficient” to execute? Why?


Questions

- Write an equivalent relational algebra expression for this query
- Are there other possible expressions?
- Do you think one would be more “efficient” to execute? Why?

```
SELECT s_name FROM student,takes,classes  
WHERE t_sid=sid AND t_cid=cid  
AND c_name='6.830'
```

Solution

SELECT s_name FROM student,takes,classes
WHERE t_sid=sid AND t_cid=cid
AND c_name='6.830'



```
⋈ (
  student,
  ⋈ (
    σ (
      classes,
      c_name = '6.830'
    ),
    takes,
    t_cid=cid
  ),
  t_sid=sid
)
```

Filtering first is probably a good idea

Filtered table is small, so do join with it and classes first

Will formalize this intuition in a few classes

IMS v CODASYL v Relational

	IMS	CODASYL	Relational
Many to many relationships without redundancy	✗	✓	✓
Declarative, non “navigational” programming	✗	✗	✓

IMS v CODASYL v Relational

	IMS	CODASYL	Relational
Many to many relationships without redundancy	✗	✓	✓
Declarative, non “navigational” programming	✗	✗	✓
Physical data independence	✗	✗	✓

Physical Independence

Can change representation of data without needing to change code

Example:

```
SELECT a.name FROM animals AS a, cages AS c  
WHERE a.cageno = c.no AND c.bldg = 32
```

- Nothing about how animals or cages tables are represented is evident
 - Could be sorted, stored in a hash table / tree, etc
 - Changing physical representation will not change SQL
- No specification of implementation
- Both CODASYL and IMS expose representation-dependent operations in their query API

IMS v CODASYL v Relational

	IMS	CODASYL	Relational
Many to many relationships without redundancy	✗	✓	✓
Declarative, non “navigational” programming	✗	✗	✓
Physical data independence	✗	✗	✓
Logical data independence	✗	✗	✓

Logical Data Independence

- What if I want to change the schema without changing the code?
- No problem if just adding a column or table
- *Views* allow us to map old schema to new schema, so old programs work
 - *Even when changing existing fields*

Key Idea: View

- View is a logical definition of a table in terms of other tables
- E.g., a view computing animals per cage

```
CREATE VIEW cage_count as  
(SELECT cageno, count(*)  
FROM animals JOIN cages ON cageno=no  
GROUP by cageno  
)
```

This view can be used just like a table in other queries

Views Example

- Suppose I want to add multiple feedtimes?
- How to support old programs?
 - Rename existing animals table to animals2
 - Create feedtimes table
 - Copy feedtime data from animals2
 - Remove feedtime column from animals2
 - Create a view called animals that is a query over animals2 and feedtimes

```
CREATE VIEW animals as (  
  SELECT name, age, species, cageno,  
    (SELECT feedtime FROM feedtimes WHERE animalid = id LIMIT 1)  
  FROM animals2  
)
```

Summary: IMS v CODASYL v Relational

	IMS	CODASYL	Relational
Many to many relationships without redundancy	✗	✓	✓
Declarative, non “navigational” programming	✗	✗	✓
Physical data independence	✗	✗	✓
Logical data independence	✗	✗	✓

Next time: Fancy SQL