


# 6.5830/1 Lecture 8



Query Optimization  
March 3, 2026

# Administration

## Important dates:

- Lab 1 writeup on March 5 in class; no electronics!
- PSet2, Project Proposals due March 10
- Quiz 1 in-class on March 12
- Lab 2 due March 19
- You should have heard from us regarding accommodations on March 5; please get in touch if not


# Join Algo Summary

| Algo                               | I/O cost                         | CPU cost                               | In Mem?  |
|------------------------------------|----------------------------------|----------------------------------------|----------|
| Nested loops                       | $ R  +  S $                      | $O(\{R\} \times \{S\})$                | R in mem |
| Nested loops                       | $\{S\} R  +  S $                 | $O(\{R\} \times \{S\})$                | No       |
| Index nested loops (R index)       | $ S  + \{S\}c \quad (c < 5)$     | $O(\{S\} \log \{R\})$                  | No       |
| Block nested loops                 | $ S  + B R  \quad (B =  S /M)$   | $O(\{R\} \times \{S\})$                | No       |
| Sort-merge                         | $ R  +  S $                      | $O(\{S\} \log \{S\})$                  | Both     |
| Hash (Hash R)                      | $ R  +  S $                      | $O(\{S\} + \{R\})$                     | R in mem |
| Blocked hash (Hash S)              | $ S  + B R  \quad (B =  S /M)$   | $O(\{S\} + B\{R\}) (*)$                | No       |
| External Sort-merge                | $3( R  +  S )$                   | $O(P \times \{S\} / P \log \{S\} / P)$ | No       |
| Simple hash ( <b>not covered</b> ) | $P( R  +  S ) \quad (P =  S /M)$ | $O(\{R\} + \{S\})$                     | No       |
| Grace hash                         | $3( R  +  S )$                   | $O(\{R\} + \{S\})$                     | No       |

Grace hash is generally a safe bet, unless memory is close to size of tables, in which case simple can be preferable

Extra cost of sorting makes sort merge unattractive unless there is a way to access tables in sorted order (e.g., a clustered index), or a need to output data in sorted order (e.g., for a subsequent ORDER BY)

# Database Internals Outline



# Query Optimization Objective

- Find the query plan of **minimum cost**
  - Many possible cost functions, as we've discussed
- Requires a way to:
  - Evaluate cost of a plan
  - Enumerate (iterate through) plan options

# Cost Estimation


- Cost Plan =  $\sum(\text{Cost Plan Operators})$
- Cost Plan Operator  $\propto$  Size of Operator Input
- Determining Size of Operator Input
  - For base tables, equal to size in storage
    - Tables with indexes may support predicate push down
  - For other operators, equal to “selectivity” x size of children
    - **Selectivity** is fraction of input size that the operator emits
    - Join selectivity defined relative to the size of the cross product

# Example (Lec 5)

SELECT \* FROM emp, dept, kids  
 WHERE sal > 10k  
 AND emp.dno = dept.dno  
 AND emp.eid = kids.eid

100 tuples/page  
 10 pages RAM  
 10 KB/page

*Selectivity*  
 $\frac{1000}{100 \times 1000} = 0.01$



$|dept| = 100 \text{ records} = 1 \text{ page} = 10 \text{ KB}$   
 $|emp| = 10K = 100 \text{ pages} = 1 \text{ MB}$   
 $|kids| = 30K = 300 \text{ pages} = 3 \text{ MB}$

Join Ordering? Why not kids / emp first?

Steps:

For each plan alternative:

1. Estimate sizes of relations
2. Estimate selectivities
3. Compute intermediate sizes
4. Evaluate cost of plan operations

Index vs scan?

5. Select best plan

Steps:

1.  Estimate sizes of relations
2. Estimate selectivities
3. Compute intermediate sizes
4. Evaluate cost of plan operations
5. Find best overall plan

# Selinger Statistics

**NCARD(R)** - "relation cardinality" - number of records in R

**TCARD(R)** - # pages R occupies

**ICARD(I)** - # keys (distinct values) in index I

**NINDX(I)** - pages occupied by index I

Min and max keys in indexes

Modern databases use much more sophisticated stats – will look at Postgres and learn about some research techniques in 2 lectures



Steps:

1. Estimate sizes of relations
2.  Estimate selectivities
3. Compute intermediate sizes
4. Evaluate cost of plan operations
5. Find best overall plan

# Selinger Selectivities

$F(\text{pred}) = \text{Selectivity of predicate} = \text{Fraction of records that a predicate does not filter}$

## Predicate types

1.  $\text{col} = \text{val}$

**NCARD(R)** - "relation cardinality" - number of records in R

**TCARD(R)** - # pages R occupies

**ICARD(I)** - # keys (distinct values) in index I

**NINDX(I)** - pages occupied by index I

Min and max keys in indexes

**Clicker** (<http://clicker.mit.edu/6.5830>)

Which is the best estimate for the selectivity of  $\text{col} = \text{val}$ ?

- A.  $1/\text{TCARD}(R)$
- B.  $\text{ICARD}(I)/\text{NCARD}(I)$
- C.  $1/\text{ICARD}(I)$
- D.  $(\text{max key} - \text{val}) / (\text{ICARD}(I))$

Steps:

1. Estimate sizes of relations
2.  Estimate selectivities
3. Compute intermediate sizes
4. Evaluate cost of plan operations
5. Find best overall plan

# Selinger Selectivities

$F(\text{pred}) = \text{Selectivity of predicate} = \text{Fraction of records that a predicate does not filter}$

## Predicate types

1.  $\text{col} = \text{val}$

$F = 1/\text{ICARD}()$  (if index available)

$F = 1/10$  otherwise



Modern DBs use fancier stats!

2.  $\text{col} > \text{val}$

$(\text{max key} - \text{value}) / (\text{max key} - \text{min key})$  (if index available)

$1/3$  otherwise

3.  $\text{col1} = \text{col2}$

$1/\text{MAX}(\text{ICARD}(\text{col1}), \text{ICARD}(\text{col2}))$  (if index available)

$1/10$  otherwise

Assumes key-foreign key join

Note a better estimate is  $1/\text{ICARD}(\text{PK table})$

**We use  $1/\text{ICARD}(\text{PK table})$  going forward**

**NCARD(R)** - "relation cardinality" - number of records in R

**TCARD(R)** - # pages R occupies

**ICARD(I)** - # keys (distinct values) in index I

**NINDX(I)** - pages occupied by index I

Min and max keys in indexes

Steps:

1. Estimate sizes of relations
2.  Estimate selectivities
3. Compute intermediate sizes
4. Evaluate cost of plan operations
5. Find best overall plan

# Complex Predicates

$F(\text{pred})$  = Selectivity of predicate = Fraction of records that a predicate does not filter

- P1 and P2

$$F(P1) \times F(P2)$$

- P1 or P2

$$1 - P(\text{neither predicate is satisfied}) =$$

$$1 - (1 - F(P1)) \times (1 - F(P2))$$

Note uniformity assumption

Steps:

1. Estimate sizes of relations
2. Estimate selectivities
3.  Compute intermediate sizes
4. Evaluate cost of plan operations
5. Find best overall plan

# Intermediate Sizes


**NCARD(R)** - "relation cardinality" - number of records in R

**TCARD(R)** - # pages R occupies

**ICARD(I)** - # keys (distinct values) in index I

**NINDX(I)** - pages occupied by index I

Min and max keys in indexes



Clicker (<http://clicker.mit.edu/6.5830>)

A)  $100 \times (10000 \times 0.1) \times 0.01 = 1000$

B)  $10000 \times 0.1 = 1000$

C)  $10000 \times 0.1 \times 0.01 = 10$

D)  $10000 \times 0.1 \times 100 = 100000$

Steps:

1. Estimate sizes of relations
2. Estimate selectivities
3.  Compute intermediate sizes
4. Evaluate cost of plan operations
5. Find best overall plan

# Intermediate Sizes


**NCARD(R)** - "relation cardinality" - number of records in R

**TCARD(R)** - # pages R occupies

**ICARD(I)** - # keys (distinct values) in index I

**NINDX(I)** - pages occupied by index I

Min and max keys in indexes



$$\begin{aligned} NCARD_d \times NCARD_e \times F_1 \times F_2 &= \\ 100 \times 10000 \times 0.1 \times 0.01 &= \\ 1000 \end{aligned}$$

Steps:

1. Estimate sizes of relations
2. Estimate selectivities
3. Compute intermediate sizes
4.  Evaluate cost of plan operations
5. Find best overall plan

Cost = pages read +  
weight x (records evaluated)

# Cost of Base Table Operations

**NCARD(R)** - "relation cardinality" - number of records in R

**TCARD(R)** - # pages R occupies

**ICARD(I)** - # keys (distinct values) in index I

**NINDX(I)** - pages occupied by index I

Min and max keys in indexes

**W**: weight of CPU operations

Heap File  
lookup

**Equality predicate with unique index:**  $1 + 1 + W$   
(Assume 1 I/O for B+Tree lookup)

B+Tree  
lookup

Predicate  
evaluation

Steps:

1. Estimate sizes of relations
2. Estimate selectivities
3. Compute intermediate sizes
4.  Evaluate cost of plan operations
5. Find best overall plan

**Cost = pages read +  
weight x (records evaluated)**

# Cost of Base Table Operations

**NCARD(R)** - "relation cardinality" - number of records in R

**TCARD(R)** - # pages R occupies

**ICARD(I)** - # keys (distinct values) in index I

**NINDX(I)** - pages occupied by index I

Min and max keys in indexes

**W:** weight of CPU operations

Heap File  
lookup

**Equality predicate with unique index:**  $1 + 1 + W$   
(Assume 1 I/O for B+Tree lookup) B+Tree  
lookup Predicate  
evaluation

**Clustered index, range w/ selectivity F**

**A:**  $F \times \text{TCARD} + W \times (\text{tuples read})$

**B:**  $F \times (\text{NINDX} + \text{NCARD}) + W \times (\text{tuples read})$

**C:**  $F \times \text{NINDX} + W \times (\text{tuples read})$

**D:**  $F \times (\text{NINDX} + \text{TCARD}) + W \times (\text{tuples read})$

Clicker (<http://clicker.mit.edu/6.5830>)

Steps:

1. Estimate sizes of relations
2. Estimate selectivities
3. Compute intermediate sizes
4.  Evaluate cost of plan operations
5. Find best overall plan

**Cost = pages read +  
weight x (records evaluated)**

# Cost of Base Table Operations

**NCARD(R)** - "relation cardinality" - number of records in R

**TCARD(R)** - # pages R occupies

**ICARD(I)** - # keys (distinct values) in index I

**NINDX(I)** - pages occupied by index I

Min and max keys in indexes

**W:** weight of CPU operations

Heap File  
lookup

**Equality predicate with unique index:**  $1 + 1 + W$   
(Assume 1 I/O for B+Tree lookup) B+Tree  
lookup Predicate  
evaluation

**Clustered index, range w/ selectivity F:**  $F \times (NINDX + TCARD) + W \times (\text{tuples read})$   
One I/O per page

**Unclustered index, range w/ selectivity F :**  $F \times (NINDX + NCARD) + W \times (\text{tuples read})$   
One I/O per record

**Seq (segment) scan:**  $TCARD + W \times (NCARD)$



Steps:

1. Estimate sizes of relations
2. Estimate selectivities
3. Compute intermediate sizes
4.  Evaluate cost of plan operations
5. Find best overall plan

# Cost of Joins

**NCARD(R)** - "relation cardinality" - number of records in R

**TCARD(R)** - # pages R occupies

**ICARD(I)** - # keys (distinct values) in index I

**W**: weight of CPU operations

## NestedLoops(A,B,pred)


$$\text{Cost(A)} + \text{NCARD(A)} \times \text{Cost(B)}$$

Outer Plan


Inner Plan

- Selinger only considers "left deep" plans, i.e., B is always a **base table** (that is, it can be loaded from storage)  $T_{\text{right}}$
- In an index on  $T_{\text{right}}$ ,  $\text{Cost(B)} = 1 + 1 + W$  (*1 I/O for B+Tree, 1 for heap*)
- If no index,  $\text{Cost(B)} = \text{TCARD}(T_{\text{right}}) + W \times \text{NCARD}(T_{\text{right}})$  (*scan the heap*)
- $\text{Cost(A)}$  is just cost of outer subtree


"right deep"



"left deep"



"bushy"



Steps:

1. Estimate sizes of relations
2. Estimate selectivities
3. Compute intermediate sizes
4.  Evaluate cost of plan operations
5. Find best overall plan

# Cost of Joins

## Merge(A,B,pred)

$$\text{Cost}(A) + \text{Cost}(B) + \text{sort cost}$$

Varies depending on whether sort is in memory or storage, and whether one or both tables are already sorted

If either table is a base table, cost is just the sequential scan cost




Steps:

1. Estimate sizes of relations
2. Estimate selectivities
3. Compute intermediate sizes
4. Evaluate cost of plan operations
5.  Find best overall plan

# Enumerating Plans

- Selinger combines several heuristics with a search over join orders
- Heuristics
  - Push down selections
  - Don't consider cross products
  - Only “left deep” plans
    - Right side of all joins is base relation
- Still have to order joins!




# Join ordering

- Suppose I have 3 tables,  $A \bowtie B \bowtie C$ 
  - Predicates between all 3 (no cross products)
- How many orderings?

|     |       |       |
|-----|-------|-------|
| ABC | A(BC) | (AB)C |
| ACB | A(CB) | (AC)B |
| BAC | B(AC) | (BA)C |
| BCA | B(CA) | (BC)A |
| CAB | C(AB) | (CA)B |
| CBA | C(BA) | (CB)A |

$n!$



This plan is not  
left deep!

Left deep plans are all of  
the form  $(\dots(((AB)C)D)E)\dots$

$n!$  left deep plans

$10! = 3.6 \text{ M}$

$15! = 1.3 \text{ T}$

Can we do  
better?

# Dynamic Programming Algorithm

- **Idea:** compute the best way to join each subplan, from smallest to largest
  - Don't need to reconsider subplans in larger plans
- For example, if the best way to join ABC is (AC)B, that will always be the best way to join ABC, whenever\* these three relations occur as a part of a subplan.

\* *Except when considering interesting orders*

# Postgres example

*explain select \* from emp join kids using (eno);*

Hash Join (cost=34730.02..132722.07 rows=3000001 width=35)

Hash Cond: (kids.eno = emp.eno)

-> Seq Scan on kids (cost=0.00..49099.01 rows=3000001 width=18)

-> Hash (cost=16370.01..16370.01 rows=1000001 width=21)

-> Seq Scan on emp (cost=0.00..16370.01 rows=1000001 width=21)

*explain select \* from dept join emp using(dno) join kids using (eno);*

Hash Join (cost=35000.04..140870.43 rows=3000001 width=39)

Hash Cond: (emp.dno = dept.dno)

-> Hash Join (cost=34730.02..132722.07 rows=3000001 width=35)

Hash Cond: (kids.eno = emp.eno)


-> Seq Scan on kids (cost=0.00..49099.01 rows=3000001 width=18)

-> Hash (cost=16370.01..16370.01 rows=1000001 width=21)

-> Seq Scan on emp (cost=0.00..16370.01 rows=1000001 width=21)

-> Hash (cost=145.01..145.01 rows=10001 width=8)

-> Seq Scan on dept (cost=0.00..145.01 rows=10001 width=8)



# Selinger Algorithm

$R \leftarrow$  set of relations to join

For  $i$  in  $\{1 \dots |R|\}$ :

for  $S$  in {all length  $i$  subsets of  $R$ }:

$\text{optcost}_S = \infty$

$\text{optjoin}_S = \emptyset$

for  $a$  in  $S$ : //  $a$  is a relation

$c_{sa} = \text{optcost}_{S-a} +$  *Cached in previous step!*

min. cost to join  $(S-a)$  to  $a$  +

min. access cost for  $a$

if  $c_{sa} < \text{optcost}_S$ :

$\text{optcost}_S = c_{sa}$

$\text{optjoin}_S = \text{optjoin}(S-a)$  joined optimally w/  $a$

# Example

## 4 Relations: ABCD

## Optjoin:

A = best way to access A

(e.g., sequential scan,

or predicate pushdown into index...)

$$B = \begin{pmatrix} & & & \\ & & & \\ & & & \\ & & & \end{pmatrix} B$$
$$C = \begin{pmatrix} & & & \\ & & & \\ & & & \\ & & & \end{pmatrix} C$$
$$D = \begin{pmatrix} & & & \\ & & & \\ & & & \\ & & & \end{pmatrix} D$$
$$\{A, B\} = AB \text{ or } BA$$
$$\{A, C\} = AC \text{ or } CA$$
$$\{B,C\} = BC \text{ or } CB$$
 $\{A, D\}$  $\{B, D\}$  $\{C, D\}$ [illegible]

## Dynamic Programming Table



# Example (con't)

| Relations | Best Plan  | Cost |
|-----------|------------|------|
| A         | Index Scan | 5    |
| B         | Seq Scan   | 15   |
| ...       |            |      |
| {A,B}     | BA         | 75   |
| {A,C}     | AC         | 12   |
| {B,C}     | CB         | 22   |
| ..        |            |      |
|           |            |      |
|           |            |      |
|           |            |      |
|           |            |      |
|           |            |      |

Already computed!

Optjoin

$\{A,B,C\} =$  remove A: compare  $A(\{B,C\})$  to  $(\{B,C\})A$   
 remove B: compare  $(\{A,C\})B$  to  $B(\{A,C\})$   
 remove C: compare  $C(\{A,B\})$  to  $(\{A,B\})C$

$\{A,C,D\} = \dots$

$\{A,B,D\} = \dots$

$\{B,C,D\} = \dots$

...

$\{A,B,C,D\} =$  remove A: compare  $A(\{B,C,D\})$  to  $(\{B,C,D\})A$   
 remove B: compare  $B(\{A,C,D\})$  to  $(\{A,C,D\})B$   
 remove C: compare  $C(\{A,B,D\})$  to  $(\{A,B,D\})C$   
 remove D: compare  $D(\{A,B,C\})$  to  $(\{A,B,C\})D$

# Complexity

- Have to enumerate all sets of size 1...n

$$\binom{n}{1} + \binom{n}{2} \dots + \binom{n}{n}$$

- Number of subsets of set of size n =  
|power set of n| =  
 $2^n$  (here, n is number of relations)

Equivalent to all binary strings of length N, where a 1 in the ith position indicates that relation i is included:

001, 010, 100, ... , 011, 111

# Complexity (cont.)

$2^n$  Subsets

How much work per subset?

Have to iterate through each element of each subset, so this at most  $n$

$n2^n$  complexity (vs  $n!$ )

$n=12 \rightarrow 49\text{K vs } 479\text{M}$

# Interesting Orders

- Some query plans produce data in sorted order –  
E.g scan over a primary index, merge-join  
– Called an *interesting order*
- Next operator may use this order – E.g. can be another merge-join
- For each subset of relations, compute multiple optimal plans, one for each interesting order
- Increases complexity by factor  $k+1$ , where  $k$ =number of interesting orders

# Summary

- Selinger Optimizer is the foundation of modern cost-based optimizers
  - Simple statistics
  - Several heuristics, e.g., left-deep
  - Dynamic programming algo for join ordering
- Easy to extend, e.g., with:
  - More sophisticated statistics
  - Fewer heuristics

